

Blind ML: An Ethical Analysis of Privacy-Preserving Threat Research

Joe Waller
DATA740: Final Project

April 15, 2026

Abstract

Blind ML is an ethical research initiative that explores how security research should be conducted when the environment under study contains sensitive, operationally dangerous, or potentially harmful user activity. Conventional abuse-detection systems generally retain raw text, message bodies, URLs, and other content-rich evidence for downstream review. These conventional approaches strengthen semantic interpretation and aid in performance enhancement efforts, but requiring a retained-data surface increases blast radius, exposes operators to harmful material, and produces an artifact that can easily become a surveillance archive. Blind ML asks whether a different posture is possible: can useful hostile-behavior analysis still be performed when raw content is removed before persistence, thereby restricting downstream learning to names-only telemetry, structural metadata, and derived scalar features?

The Moltbook Observer Project is the implementation testbed for that question. Its architecture combines edge obfuscation, schema-bound packet families, staged simulation and collection lanes, fail-closed posture controls, hardened runtime boundaries, and local analysis workflows. The strongest ethical claim of the project is not that metadata-only learning eliminates platform-moderation risk. The strongest claim is that privacy, safety, and governance can be treated as binding requirements inside an ML workflow rather than as a compliance language attached after the system is already built. Viewed in its intended mature form, Blind ML becomes a governed multi-regime research program in which synthetic data, canary data, live-adjacent data, and reviewed release views all contribute distinct kinds of evidence without breaking the underlying privacy contract.

This report evaluates that mature design as an ethical data-science case. It argues that the project is compelling not because it eliminates every uncertainty in hostile-behavior detection, but because it demonstrates a disciplined way to study such behavior without normalizing unsafe retention.

1 Introduction

Much of contemporary platform-safety work assumes that raw content is the natural substrate of detection. If an investigator wants to study phishing, harassment, prompt injection, social engineering, or malicious coordination, the default pipeline captures the original text and treats richer semantic access as a prerequisite for meaningful modeling. That assumption has practical advantages, but it also creates ethical and operational burdens. Retained raw content increases the number of people and operational surfaces exposed to harmful material and amplifies the consequences of loss or misuse of data. In an agentic environment, where automation can collect, process, and act on information quickly and at scale, these burdens become more consequential.

These governance concerns are prominent in increasingly-complex agent-systems discussions, where security, interoperability, and operational control are being treated as first-class policy questions rather than background implementation details [1, 2].

Project repository: <https://github.com/joediggidy/calamum-moltbook-observer/>. Tooling disclosure: This project was developed using a repository governance harness and an AI-assisted coding/curation workflow engineered by Joe Waller and enabled via an OpenAI API. These tools are used to support development and documentation, including summarization, and do not replace the author's judgment. Names such as CodeSentinel, ORACL, and Calamum appear as internal identifiers for governance, agent assistance, and ethical-security harnesses, respectively.

Blind ML begins with a different assumption. Instead of asking how much content a security-research stack can safely retain, it asks how much useful inference can survive after semantic payloads have already been stripped away. The guiding hypothesis is not that metadata is sufficient for every security task; the hypothesis is that structural, temporal, and interaction-level signals may remain useful enough to support a meaningful class of hostile-behavior research while dramatically reducing the privacy and exposure costs associated with content-rich pipelines.

Previous work on metadata-based abuse detection and privacy-preserving moderation supports the plausibility of this direction by showing that non-semantic or reduced-semantic features can still contribute meaningfully to safety-relevant modeling tasks [3, 5, 4].

That hypothesis makes the Moltbook Observer strong as an ethical investigation, because the ethical question is inseparable from the technical design. The central issue is not simply whether a model can be trained, but whether a model can be trained under the honor privacy, proportionality, operator safety, and honest reporting constraints proposed by this framework. The observer is therefore best understood not merely as a software artifact but as a case study in ethical security by design.

2 Role, Sector, and Practical Problem

The role represented by this project can best be described as a hybrid of an ethical security researcher, a data guardian, and a defensive system designer. The sector is cyber-security and platform integrity, with a direct secondary connection to applied machine learning and responsible data governance. These roles matter because they define the standard against which the work should be judged. A commercial moderation product might optimize first for throughput, coverage, or intervention speed. An ethical security-research project must also justify how its measurement practices affect the people and systems it studies. This role-and-sector framing is also reflected in the project’s formal alignment assessment.

The practical problem is straightforward to state and difficult to solve well. Online platforms need ways to study suspicious or hostile behavior. Researchers need enough signal to identify patterns worth investigating. Operators need tools that are reproducible, controlled, and auditable. At the same time, platform subjects retain privacy interests even when their behavior is public-facing or risky, and operators should not be casually exposed to high-risk content when a safer design is possible. The project therefore confronts a familiar tension between security utility and private-data restraint, but it does so in a context where governance failures have immediate technical consequences.

The motivating question can be stated plainly: how can a research system study hostile activity without turning itself into a second site of harm? Harm in this context is not limited to breach risk. It includes operator exposure, quiet over-retention, mission creep, misleading reporting, and the gradual normalization of expansive collection simply because richer data is easier to model. Blind ML responds by treating bounded observation as part of the scientific method. That decision makes the learning problem harder, but it also makes the ethical stakes explicit.

3 Blind ML and the Observer Testbed

Blind ML is best defined as a *privacy-preserving machine-learning posture* rather than a single algorithm. Its core claim is that security-relevant research can be organized around obfuscated, schema-bound telemetry rather than retained message bodies. In the Calamum system, downstream analysis sees timestamps, event types, counts, flags, hashed identifiers, prompt-risk helper indicators, and derived scalar summaries such as complexity or code density. It does not depend on storing the underlying body text as the primary modeling substrate.

The Calamum Moltbook Observer turns that posture into an executable testbed. The public project overview and methodology manual describe a system with five clear commitments: synthetic-first validation, names-only telemetry, strict separation of scientific telemetry from diagnostics and safety governance, fail-closed enforcement, and reproducible local analysis. Privacy, accountability, and operator restraint are treated not only as first-class goals, but have the ability to determine if the project is allowed to proceed at all. The relevant project evidence surfaces are the public project overview and the methodology manual.

The project is especially interesting when understood through its architecture. The dataset authority layer does not simply pass downstream tooling whatever columns happen to appear in a raw retained table. Instead, reviewed dataset views are released with explicit field roles already classified. Identity, provenance, governance, and modeling features are separated before train, evaluate, and score stages consume the data. Ethically, that mechanism matters because it prevents downstream convenience from redefining what the data is supposed to mean. Interpretation and admissibility remain tied to reviewed release packets rather than ad hoc feature selection performed late in the pipeline.




Figure 1: Governed Blind ML flow. Raw platform activity is reduced at the edge, converted into names-only telemetry, and released downstream only through reviewed dataset views and public reporting surfaces.

The strongest version of this paper is not a vague claim that the project is ethically thoughtful or operationally complete. It is a concrete description of what the fully realized experiment program is intended to look like. In its ideal state, Blind ML is not one model run against one convenient file. It is a controlled comparison across multiple data regimes, model families, thresholds, and reporting surfaces, all governed by the same names-only privacy contract. While the underlying premises and operational tooling have been thoroughly tested shown to be effective, the more robust operational features such as granular feature selection and dataset merging have been deferred to version 2 development.

3.1 Research Questions

The mature experiment lane is organized around four research questions. First, can structural metadata alone distinguish ordinary behavior from higher-risk behavior at a level that is useful for research triage? Second, can that usefulness be preserved across different data regimes, including synthetic, canary, live-adjacent, and reviewed honeypot slices? Third, can a conservative operating threshold be maintained so that the detector remains practically usable rather than simply sensitive? Fourth, can the governance architecture preserve truthful claims, reproducibility, and bounded retention while the experiment grows in complexity?

Taken together, these questions define success in a way that is distinct from ordinary leaderboard-style ML. The ideal outcome is not maximal predictive power at any cost. The ideal outcome is disciplined utility under a privacy-preserving observation contract.

3.2 Dataset Regimes

The mature experiment program relies on multiple dataset regimes because each one serves a different empirical and ethical role.

This regime-based design is important because it prevents the project from collapsing every research question into a single dataset and a single number. Synthetic data answers one class of question, benign canary data answers another, and real but names-only telemetry answers another still. The ethical benefit of this staged structure is that realism is introduced gradually and deliberately rather than treated as harmless by default.

3.3 Model Families and Comparison Strategy

Instead of treating one library call as the whole scientific contribution, the Observer’s experiment program compares multiple families of models. Heuristic baselines remain useful because they expose obvious structural rules and provide a floor for interpretation. Isolation Forest and related anomaly methods matter because the project can learn from unlabeled or weakly labeled regimes where conservative triage is the primary goal. Supervised models such as Logistic Regression and Random Forest become important once reviewed labels or

Dataset regime	Primary purpose	Ethical function	Ideal-state use
Synthetic TV-0 to TV-3 corpus	establish controlled truth conditions	supports low-risk validation before higher-risk observation	baseline supervised and unsupervised model development
Canary benign corpus	calibrate conservative false-positive posture	honors the precautionary principle by testing on safer traffic first	threshold tuning and benign-score distribution analysis
Live-adjacent names-only corpus	capture realistic operational structure	grounds the project in real telemetry without expanding raw-content retention	drift analysis, score distribution checks, robustness testing
Provenance-labeled companion view	support bounded reviewed comparisons	preserves label honesty by declaring what labels actually mean	limited supervised experiments with explicit interpretation boundaries
Reviewed honeypot slice	concentrate higher-risk but controlled observations	creates a tighter link between measurement and reviewed risk posture	focused high-risk evaluation and case-study analysis

Table 1: Dataset regimes in the intended mature Blind ML experiment program.

acceptable proxy labels exist, because they test whether the metadata-only regime supports discrimination beyond outlier ranking alone.

This comparison strategy is ethically relevant in two ways. First, it avoids overstating the evidence from a single favorite method. Second, it respects the fact that privacy-preserving ML should be judged partly by how robust its conclusions are to methodological variation. An ethical research program should not let the boundary conditions of one algorithm become the only story told about a privacy-constrained feature space.

3.4 Evaluation Criteria

The ideal-state evaluation lane reports precision, recall, F1-score, false-positive rate, and threshold behavior wherever labels permit. Where labels are partial or intentionally bounded, the most relevant measures are ranking quality, threshold conservatism, and cross-regime stability. This is especially important because a privacy-preserving detector becomes ethically weaker if it floods benign activity with alarms and then asks human reviewers to absorb the cost.

Success is therefore defined in layers. At the first layer, the project must show that the metadata-only feature space is non-trivial. At the next, it must show that conservative thresholds can be set and reproduced. At the strongest layer, it must show that multiple models trained on admissible dataset views converge on results that are interpretable, bounded, and honest about their uncertainty. Those layers together form the empirical counterpart to the ethical thesis.

4 Stakeholders and Ethical Matrix

The ethical matrix for this project is strongest when it stays close to concrete stakeholders rather than abstract principles alone. Four stakeholder groups matter most: platform subjects whose behavior is being measured, researchers and operators who run the system, platform owners or governance actors who authorize or review the work, and public readers such as graders, peer reviewers, or auditors who rely on the final report.

For platform subjects, the major ethical question is whether the system preserves enough restraint to justify observation at all. Edge obfuscation and names-only persistence matter because they narrow the retained-data surface before downstream analysis occurs. That does not eliminate privacy risk; metadata can still reveal patterns of timing, activity intensity, and interaction structure. But it changes the nature of the risk from broad semantic retention to a more bounded inference problem, which is ethically preferable if the research question can still be answered.

Stakeholder group	Primary good sought	Main ethical risk	Design choices that matter most
Platform subjects	privacy, proportionality, bounded observation	metadata inference, over-collection, misclassification	edge obfuscation, names-only telemetry, staged escalation
Researchers and operators	safety, reproducibility, interpretability	exposure to harmful content, unsafe runtime shortcuts, overconfidence	synthetic-first validation, hardened runtime, fail-closed controls
Platform owners and governance actors	accountable monitoring, auditable controls	mission creep, opaque evidence chains, weak transition discipline	posture gates, selector-based dataset governance, report lineage
Readers, reviewers, and auditors	truthful claims, clear evidence boundaries	polished but misleading reporting, hidden gaps between local and public artifacts	citation packs, explicit packet routing, honest claim boundaries

Table 2: Primary stakeholder groups and ethical trade-offs in the Blind ML design.

For researchers and operators, safety is part of the method. A research environment that quietly exposes staff to harmful content, expands permissions for convenience, or rewards unsafe shortcuts is not ethically rescued by a good model score. This is why the hardened container boundary, non-root execution, and fail-closed posture checks matter so much in the paper. They are not side features; they are part of the system’s moral scaffolding.

For platform owners and governance actors, the project’s strongest contribution is that it makes authority explicit. Dataset selection, mode transitions, and report publication are not supposed to be implied by scattered artifacts. They are governed by named control-plane surfaces and explicit release or selector decisions. In the ideal state of the project, that makes it harder for downstream analysis to inherit un-reviewed data assumptions by accident.

For readers and reviewers, the ethical obligation is honesty. The report should help the reader see what is real, what is partial, and what remains open. That is why the project’s ethical matrix ties privacy, accountability, operator safety, traceability, and honest reporting together. Utility without clear boundaries would weaken the project rather than strengthen it.

5 Data Access, Selection, Curation, and Privacy

The project’s data methodology manual provides the clearest evidence for its privacy posture. The methodology manual defines names-only packet families for content items, interaction events, runtime envelopes, baseline telemetry, evidence packets, and public report packets. Across these families, the retained fields focus on counts, flags, hashes, derived scalars, event metadata, and run linkage. This matters because ethical claims about privacy must be tied to what the system actually retains, not to aspirational statements about being careful.

The project is strongest where it treats the retention boundary as a first-class contract. The observer retains structure and behavior signals, not raw semantic payloads, as the primary research substrate. That does not make the data ethically neutral. Hashed identifiers, timing patterns, and metadata correlations can still support inference. But it creates a principled narrower surface that is easier to justify than ordinary content-centric logging. The public project overview and methodology manual make this boundary explicit in the repository documentation.

Selection and curation are equally important. The data story becomes most coherent when it recognizes that not every dataset view should be treated as morally or analytically interchangeable. Synthetic corpora offer explicit truth conditions. Benign canary data supports conservative threshold calibration. Live-adjacent names-only telemetry offers realism but often weaker labels. Reviewed honeypot slices offer concentrated structure under higher scrutiny. In the ideal-state architecture, these are not merely separate files. They are separately governed releases with explicit field-role definitions and explicit interpretation boundaries.

That curation model is ethically preferable to a looser pipeline because it reduces the chance that downstream code will quietly treat provenance, identity, or governance metadata as if it were modeling signal. The mature system releases a reviewed feature view for model consumers rather than a mixed-authority raw table. That design preserves both privacy and analytical honesty. It also gives the paper a stronger answer to the curation question: the project does not simply minimize data quantity, it also minimizes ambiguity about what each retained field is allowed to mean.

The label question deserves special emphasis. Provenance labels, review flags, and source-lane markers can support bounded experiments, but they are not the same as semantic ground truth about hostile content. One of the project’s ethical strengths is its willingness to preserve that distinction rather than smoothing it over for narrative convenience.

6 Governance Structures and Defense in Depth

Governance in this project is implemented as layered technical control. At the methodology level, the system commits to synthetic-first validation, names-only persistence, strict separation of telemetry from governance signals, and reproducible local analysis. At the runtime level, it relies on multiple protective layers: obfuscation before persistence, a hardened Docker boundary, and a watchdog runtime that monitors for fail signatures and can terminate unsafe states. At the data-authority level, selector and manifest logic distinguish between ordinary local artifacts and reviewed dataset surfaces intended for modeling or reporting.

This layered structure matters because it makes governance inspectable. A privacy-preserving ML paper is weaker if the privacy claim lives only in prose while the operational stack remains permissive by default. By contrast, the current design points to concrete controls. The container hardening report demonstrates read-only filesystem enforcement, non-root execution as UID 10001, and removal of unnecessary binaries such as `ping`. The methodology manual identifies the runtime boundary and the expected fail-closed posture. The public project overview distinguishes public documentation from local evidence and makes the reporting boundary explicit. These claims are documented in the container hardening verification report, the data methodology manual, and the public project overview.

The mature version of this governance model is especially compelling because it treats the report lane as part of the ethics contract. Public-facing reporting surfaces are not supposed to invent certainty that the local authority layer does not support. Reporting discipline is therefore not secondary to the project’s ethics, it is one of the ways the project enforces them. When the system declines to overstate labels, thresholds, or evaluation coverage, it is aligning the narrative layer with the truth conditions of the evidence layer.

7 Evidence, Claim Boundaries, and Limitations

The current project materials already support several claims with substantial confidence. They support the claim that privacy-preserving telemetry is operationally defined; the data methodology manual describes concrete packet families and invariants. They support the claim that runtime restraint is implemented rather than merely promised; the hardening report and related runtime surfaces document actual containment controls. They support the claim that reporting and dataset authority matter; the project distinguishes between reader-facing public packets, local canonical artifacts, and reviewed dataset or citation surfaces.

Two retained figures from the current observer analysis lane are especially useful for this paper because they make the project’s claim boundaries visible. They do not prove full hostile-behavior classification. They do show that the current metadata-only anomaly lane has a structured score landscape and a legible threshold posture on benign canary traffic, and that each collection mode workflow analysis produces data useful to downstream processing layers in building an ever-richer and more intelligent model.




Figure 2: Benign canary anomaly-score distribution on a log-count scale. The broad, structured score landscape indicates that the metadata-only feature regime supports a non-trivial baseline rather than collapsing into a flat score mass.




Figure 3: Threshold impact analysis for the benign canary calibration run. The selected operating point isolates a bounded flagged region while maintaining an approximately 1% benign false-positive rate.

What the repository does not yet fully prove on its own is the strongest possible form of the ML argument. The ideal-state experiment still needs complete cross-regime model comparison, stable released feature views for all modeling lanes, and evaluation artifacts that can support broader performance claims without stretching the meaning of partial labels. This is not a weakness in the ethical framing of the paper. It is part of what makes the project ethically credible.

The limitations are therefore real and important. Metadata-only learning imposes a ceiling on what the system can know. Content-free features may preserve timing, structure, and interaction patterns, but they can miss semantic nuance that matters for certain categories of threat or abuse. Data minimization also does not eliminate privacy risk. Metadata can still reveal activity rhythms, coordination, and other behavioral signatures. Finally, bounded label surfaces and reviewed slices remain bounded. They must not be converted into stronger semantic conclusions than they warrant.

These limitations do not weaken the ethical value of the project. They strengthen it because the project insists that privacy, minimization, and governance remain binding even when they make the empirical problem harder. The premise that drives this project is not incumbent upon making ethical security easy, the goal is to define a surface that prioritizes personal data rights and agent-focused operational safety as an ethically-minded front line of defense.

8 Recommendations for Mature Deployment and Study

The first recommendation is to keep dataset release authority close to the Librarian layer so downstream consumers receive reviewed views rather than mixed-authority raw tables. The second is to compare multiple model families and multiple dataset regimes so the privacy-preserving claim is not tied to one fragile method. The third is to publish threshold rationale, dataset-view identity, and claim boundaries alongside every major result packet. The fourth is to treat human review and governance gates as permanent features of the design rather than as temporary scaffolding to be discarded once the modeling stack becomes more capable.

Together, these recommendations point toward an ethical research program that matures without abandoning its founding constraints. The project is most convincing when it does not trade away privacy or honesty in the name of apparent progress.

9 Conclusion

Calamum Moltbook Observer is a strong DATA740 case because it turns ethics into system design. Blind ML, in this project, is not simply a preference for privacy over performance. It is an attempt to show that privacy, operator safety, disciplined curation, and honest reporting can be built into the architecture of a security-research workflow from the start. Edge obfuscation, names-only telemetry, staged data regimes, hardened runtime controls, selector-based dataset authority, and explicit report boundaries all contribute to a coherent ethical method.

Viewed in its ideal experimental state, the project offers a compelling answer to the course’s core questions. It identifies a socially meaningful technical problem, grounds that problem in a real sector and stakeholder structure, and responds with a design whose ethical meaning is expressed through concrete choices about data, governance, and evidence. Its broader contribution is therefore not that it eliminates every uncertainty in hostile-behavior detection. Its contribution is that it demonstrates a more defensible way to ask those questions in the first place.

Blind ML is worth taking seriously as an ethical data-science proposal because it argues that useful security research does not have to begin by normalizing unsafe retention. It can begin with constraint, stay accountable as it scales, and still produce evidence worthy of study.

Project Evidence Surfaces

The internal project documents most directly supporting this paper are the public project overview, the data methodology and packet contract, the formal alignment assessment, the Blind ML ethical matrix, and the

container hardening verification report. These documents are treated as project evidence surfaces rather than external bibliography items.

References

- [1] Announcing the “AI Agent Standards Initiative” for Interoperable and Secure Innovation. 2026. NIST, February 18. <https://www.nist.gov/news-events/news/2026/02/announcing-ai-agent-standards-initiative-interoperable-and-secure>.
- [2] CAISI issues request for information about securing AI agent systems. 2026. NIST, January 12. <https://www.nist.gov/news-events/news/2026/01/caisi-issues-request-information-about-securing-ai-agent-systems>.
- [3] M. Casavantes, M. E. Aragón, L. C. González, and M. Montes-Y-Gómez. Leveraging posts’ and authors’ metadata to spot several forms of abusive comments in Twitter. *Journal of Intelligent Information Systems*, 61(2):519–539, 2023. <https://doi.org/10.1007/s10844-023-00779-z>.
- [4] P. Leonidou, N. Kourtellis, N. Salamanos, and M. Sirivianos. Privacy-preserving online content moderation: A federated learning use case. *arXiv preprint arXiv:2209.11843*, 2022. <https://arxiv.org/abs/2209.11843>.
- [5] Z. Zhou and T. Islam. Power of Social Media Meta Data: Enhancing Cyberbullying Detection. In *Advanced Sciences and Technologies for Security Applications*, pages 339–361, 2025. https://doi.org/10.1007/978-3-031-82031-1_19.